

# Reducing Memory Access Latencies using Data Compression in Sparse, Iterative Linear Solvers

**Neil Lindquist**

All-College Thesis Defense

April 16th, 2019

# Goal

- Sparse linear systems used by many scientific computations
- Problems can be large, with over a million variables
- Arithmetic is faster than fetching data from memory

# Solver Description

- Preconditioned Conjugate Gradient was used for tests
    - Preconditioned with a 3 level multigrid
    - Symmetric Gauss-Seidel step smoother
  - Matrix store in CSR format
    - Stores the column index and value for each nonzero entry
- 3 compressible data structures
- Vector Values
  - Matrix Indices
  - Matrix Values

## Main Data Access Pattern

```
for row in rows do  
  for nonzero entry in row do  
    LOAD entry's value  
    LOAD entry's column index  
    LOAD vector value for column index  
  end for  
  WRITE vector value for row  
end for
```

- need random vector reads
- need vector writes
- need both forward and backward iteration of matrix rows

## Analytical Performance Model

- System of equations that assumes no processor-level parallelism
- Solving for what conditions outperform the baseline gives:

$$\text{totalDecode} < 32.5375 - 0.037037 \cdot \text{vectEncode}$$

$$\begin{aligned} \text{matBytes} < & 12.9664 - 0.398506 \cdot \text{totalDecode} \\ & - 0.0147595 \cdot \text{vectEncode} \end{aligned}$$


$$\begin{aligned} \text{vectBytes} < & 7.9998 + 8.27835 \cdot (12 - \text{matBytes}) \\ & - 3.29897 \cdot \text{totalDecode} - 0.122184 \cdot \text{vectEncode} \end{aligned}$$

## Simulation Performance Model

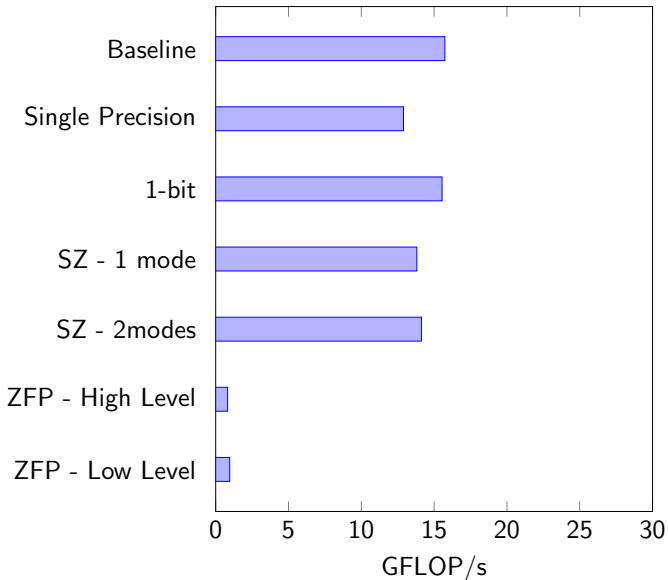
- Some processor level parallelism
- Conditions for outperforming the baseline:

Bytes	Matrix Index Decode	Matrix Value Decode	Vector Encode and Decode
1	66	41	$4.75 \geq 1.75 \cdot \text{decode} + \text{encode}$
2	51	14	$4.75 \geq 1.75 \cdot \text{decode} + \text{encode}$
3	66	40	$2 \geq 2 \cdot \text{decode} + \text{encode}$
4	0	0	$0 = \text{decode} = \text{encode}$
5	-	37	$4.75 \geq 1.75 \cdot \text{decode} + \text{encode}$
6	-	9	Not Possible
7	-	35	$0 = \text{decode} = \text{encode}$
8	-	0	$0 = \text{decode} = \text{encode}$

## Vector Compression

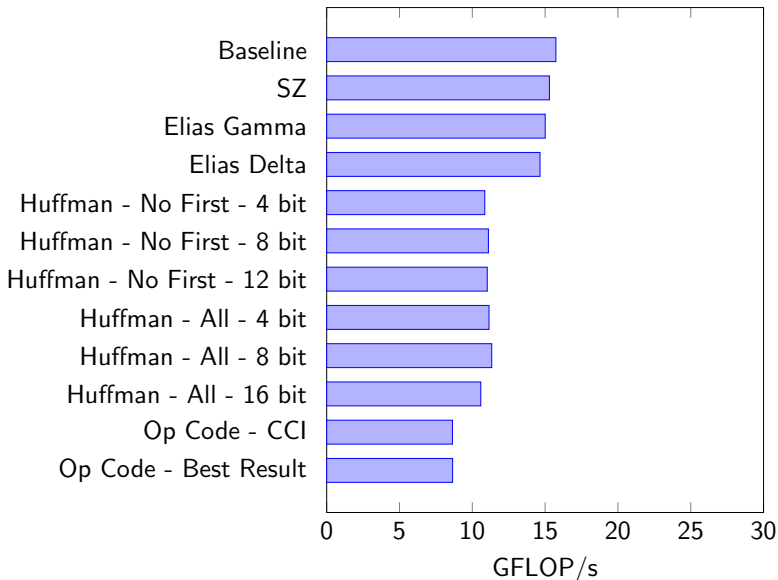


## Matrix Value Compression







## Matrix Index Compression



## Matrix Value and Index Compression



# Vector and Matrix Compression



# Successful Compression Methods

- Mixed Floating Point Precision
- SZ Compression
- Elias Gamma Coding

## Mixed Floating Point Precision

- Single Precision takes half the storage space
- But drops from 15-17 significant digits to 6-9 digits
- Certain vectors can be lower precision without slowing convergence

## Squeeze “SZ” Compression

- Stores a key for the predictor function with the best accuracy
- If can't be predicted within a minimum accuracy, sorted explicitly in a second list
- Available prediction functions are chosen based on the type of data

# Elias Gamma Coding

- Positive integers
- Stores number of bits needed then the data
- Very effective for small integers
  - “1” takes 1 bit
- Elias Delta Coding is similar, but Gamma codes the length

# Conclusion

- Iterative linear solvers are memory access bound
- Compressing key data structures provided an 84% increase in performance